

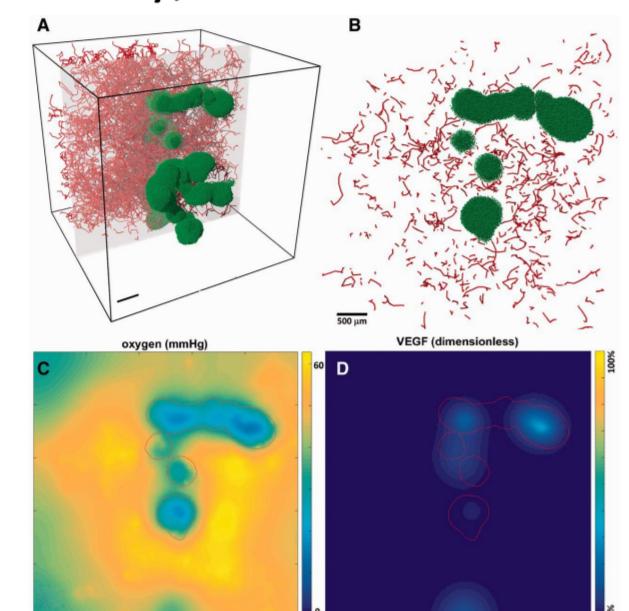
## MOTIVATION AND GOAL

### MOTIVATION

- BioFVM efficiently models the systems of PDE's for multicellular diffusion, decay, uptake, and release of multiple substrates in 2D and 3D environments.
- BioFVM is parallelized by OpenMP for efficient transport solving on CPU, but could see efficiency gains running on GPUs.
- Scientific possibilities of accelerated transport solving include higher simulation throughput, finer resolution, higher accuracy, and new (previously infeasible) 3D problems.

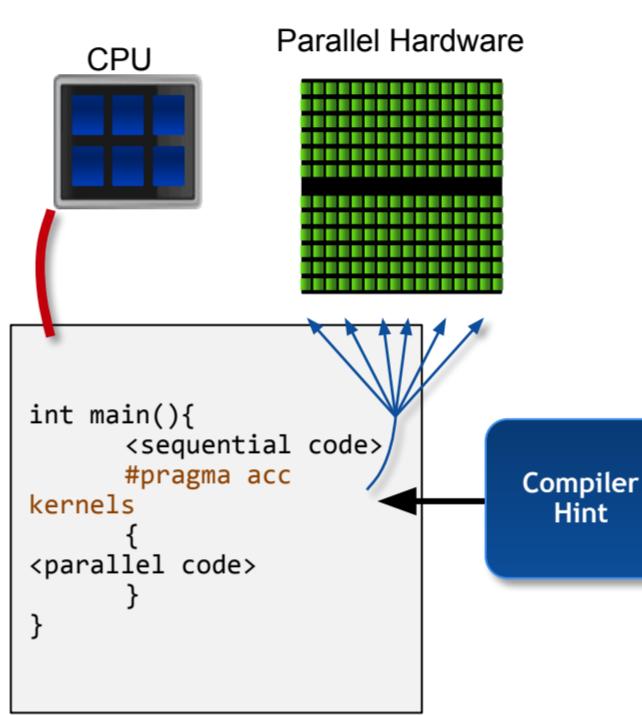
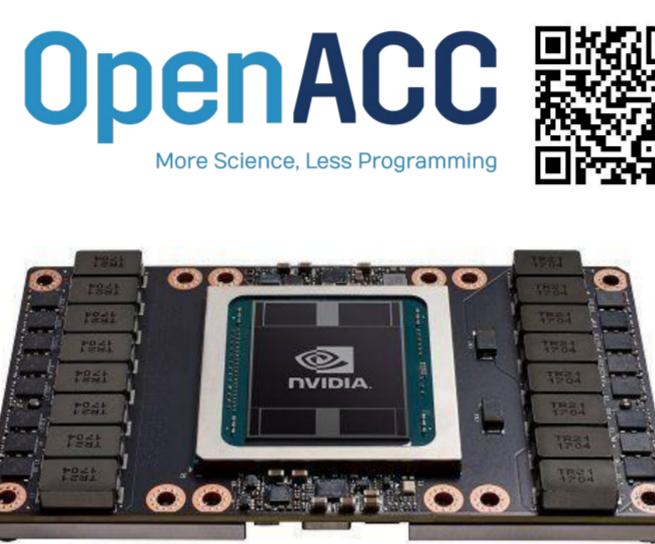
### GOAL

- Create a GPU accelerated version of BioFVM using OpenACC capable of running on the range of multicore CPU to supercomputer



## USING OpenACC DIRECTIVES

- In order to maintain portability of many different hardware architectures, and to ease the development process of working on a pre-existing code, we decided to use OpenACC to parallelize the code
- OpenACC [3], a directive-based parallel programming model for accelerators



## PROJECT ROADMAP

- BioFVM is tens of thousands lines of code, and for 2D problems the hot spot is x and y diffusion, each are a short but very complex segment of calculations
- Becoming even slightly familiar with the advanced Biology and PDE's took lots of time and help
- Tracking the all the data in the form of pointers, C++ class members, and others was fascinating to watch

$$\frac{\partial \rho}{\partial t} = \vec{D} \nabla^2 \vec{\rho} - \vec{\lambda} \vec{\rho} + \vec{S}(\vec{\rho}^* - \vec{\rho}) - \vec{U} \vec{\rho}$$

sources and uptake by cells

$$+ \sum_{cells,k} 1_k(\vec{x}) [\vec{S}_k(\vec{\rho}_k^* - \vec{\rho}_k) - \vec{U}_k \vec{\rho}_k]$$

in  $\Omega$

May

- Our first major roadblock was with C++ STL Vectors
- This data structure obscures the data from us and makes it difficult to manage CPU-GPU data
- We replaced vectors with C-style arrays when possible, and in some cases we could simply use data() to get the underlying pointer for our OpenACC data constructs

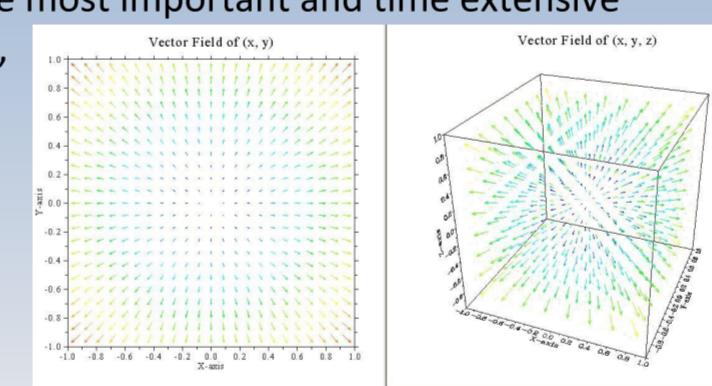
```
#pragma acc enter data
create(this->temp_thomas_cy[0:bin_thomas_cy][0:0])
for (int i = 0; i < bin_thomas_cy; i++) {
    int size = thomas_cy[i].size();
    temp_thomas_cy[i] = thomas_cy[i].data();
    #pragma acc enter data
    copyin(this->temp_thomas_cy[i:1][:size])
}

• OpenACC is directed to "attach" the temp_thomas_cy 2D vector to the C++ object calling this function, allowing correct pointer manipulation in GPU address space
```

Very expensive operation in an implementation supposed to have dramatic speedups

- Minimizing the number of times this transfer is called is the key to having a good balance of communication and computation

- Moving from 2D → 3D requires new functions to be ported to Device and introduce new z-diffusion on GPU
- A user will be able to control with ease whether BioFVM should be run parallelized with OpenACC on GPU, or for smaller domain problems/when there's no GPU available run on CPU with the existing OpenMP accelerated code
- Solving diffusion is the most important and time extensive piece of computation, but more background work done on the Device equates to less time wasted on communication



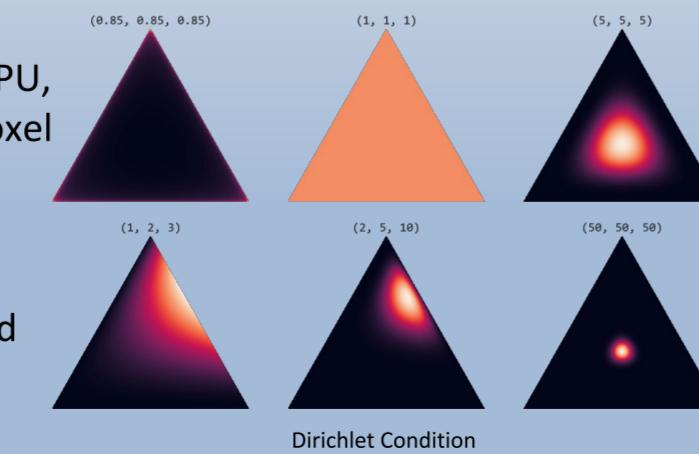
Next Steps

June

- Converting the gcc compiled code to an intelligent PGI compiler that allowed for profiling with PGPROF
- Maintain accuracy during transition between compilers while exploring the new features of PGI
- Discovered the errors that helped guide the work including the ambiguous "Illegal Address during Kernel Execution"
- Explored PGI feature Managed Memory which bridges the gap between Host and Device memory, but not used as it's not standard across compilers

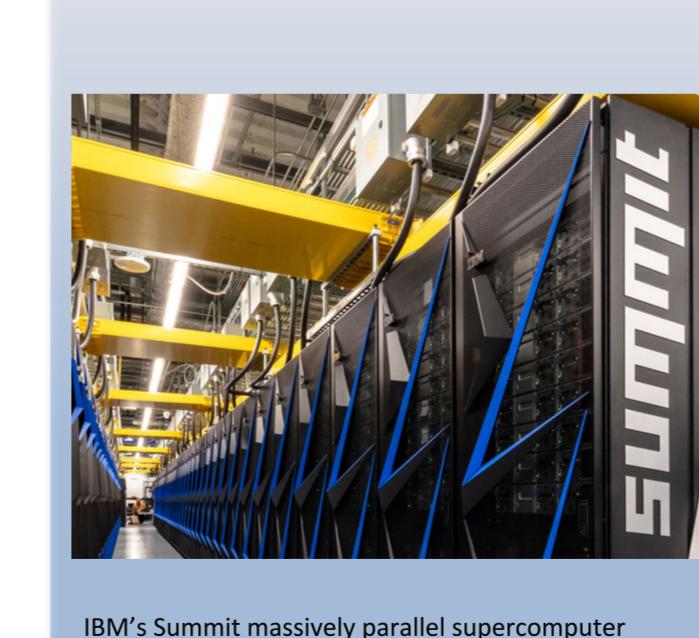
August

- Making correct copies of the rest of the object data to allow for more time computing on the GPU and less time updating between Host and Device versions of p\_density\_vector
- Constructing parallelized version of functions used in accelerated region, including AXPY and NAXPY
- Converting mesh dirichlet condition application to GPU, along with porting each voxel dirichlet data
- x-diffusion, y-diffusion, and dirichlet\_conditions now running uninterrupted on GPU



Future Work

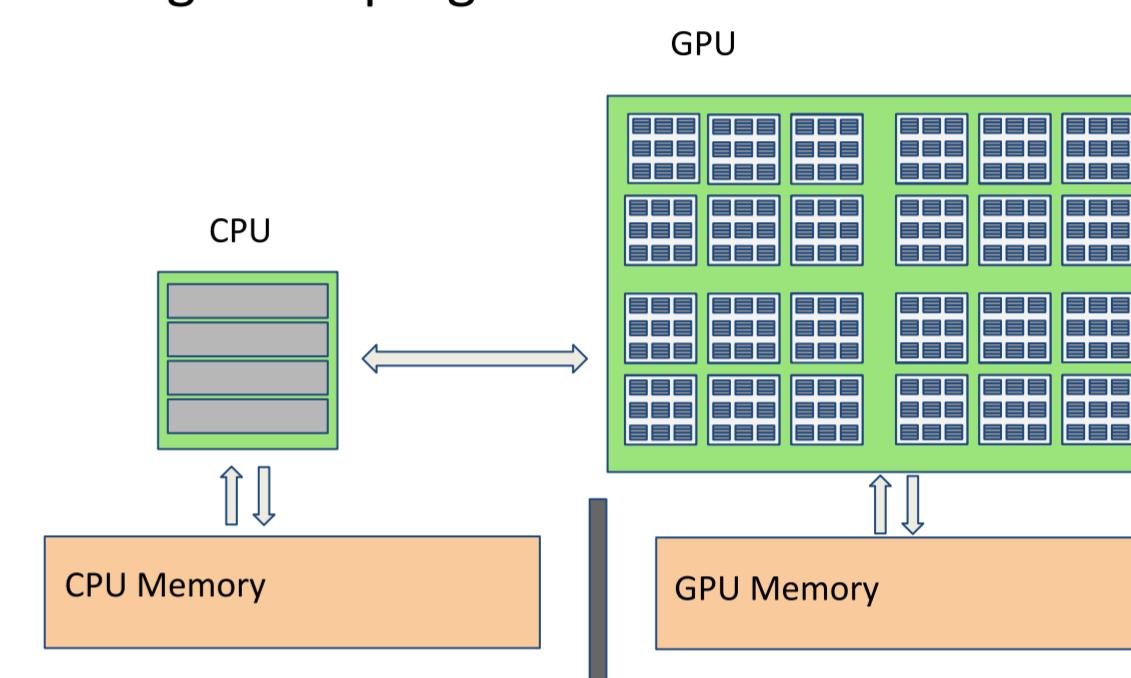
- Full computation on GPU
- One transfer of data to the Device at the beginning of the simulation, and one transfer back to Host at the end
- Full documentation with examples and instructions how to follow the data the user cannot see
- Once PGI Managed memory is a standard feature to all of OpenACC, reassessing the pros and cons of a Managed memory model (very simple, yet slower than manual)



## HANDLING DATA

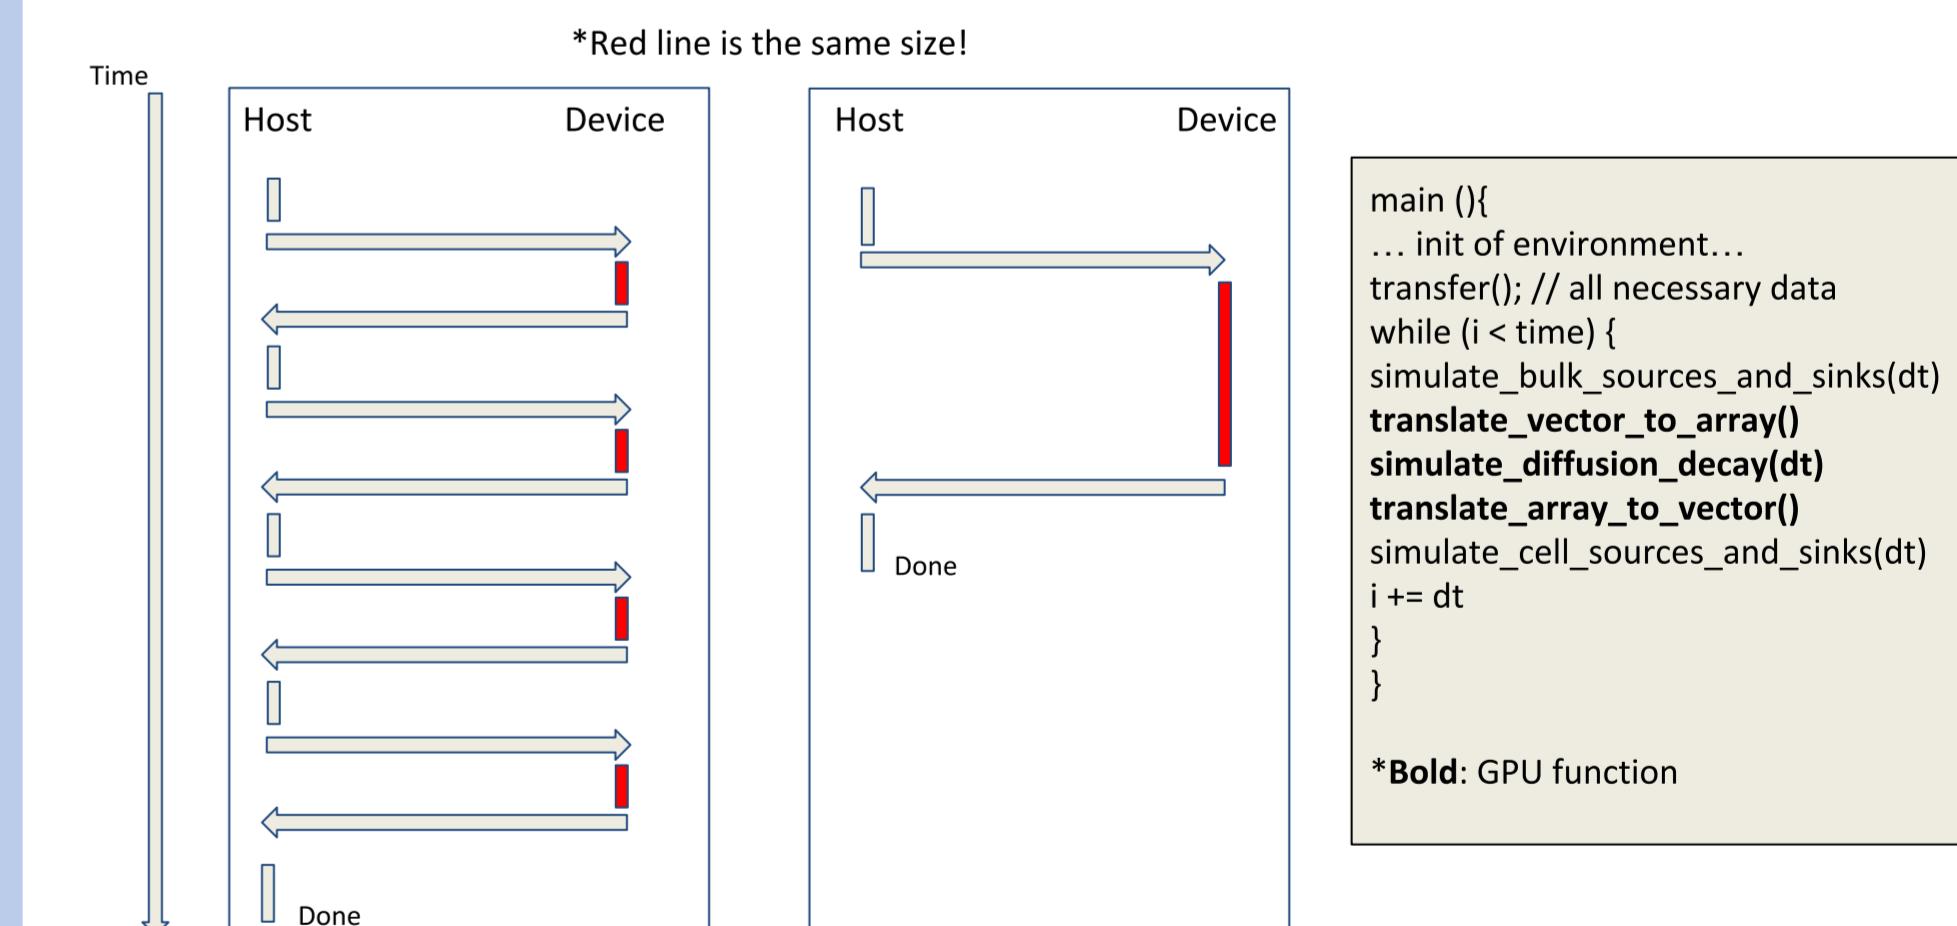
- Data must be explicitly transferred between Host and Device.
- Vector of vectors **p\_density\_vectors** holds the values of the PDE's.
  - Translation: Converting data from CPU data type (vectors and 2D vectors) to and from GPU data type (arrays and 2D arrays) while maintaining accuracy
  - Transfer: Iterating through data, dereferencing the pointers and using OpenACC data migration pragmas to transfer

Translating and transferring requires communication between the Host (CPU) and the Device (GPU) and is necessary but extremely expensive so the goal is to minimize communication!



## Results

Current results are actually slightly slower than the CPU version! But this is completely expected and correct and here's why:  
 Until every part of the parallel region (main work loop) is available on GPU, there is slow down from communication between translating and updating the two copies of p\_density\_vector.



## References & Acknowledgements

- Ahmadreza Ghaffarizadeh, Samuel H. Friedman, Paul Macklin, BioFVM: an efficient, parallelized diffusive transport solver for 3-D biological simulations, Bioinformatics, Volume 32, Issue 8, 15 April 2016, Pages 1256–1258
- The Portland Group, PGI Compilers and Tools User's Guide, [www.pgroup.com](http://www.pgroup.com) April 2018.
- <https://www.openacc.org/>

I would like to thank Mathew Colgrove from NVIDIA for his help and examples!



- Discovered the errors that helped guide the work including the ambiguous "Illegal Address during Kernel Execution"
- Explored PGI feature Managed Memory which bridges the gap between Host and Device memory, but not used as it's not standard across compilers